

Alexandre Piché

Curriculum Vitae

☎ (+1) 514 943 9865
✉ alexandrepiche@gmail.com
🌐 alexpiche.github.io

Education

2024 Université de Montréal.

Ph.D. in Computer Science, Mila.

Thesis: Searching for Q^* .

2017 McGill University.

M.Sc. Mathematics and Statistics.

Thesis: Bayesian nonparametric modelling of heterogeneous groups of censored data.

Work Experience

2022-Now ServiceNow Research, Montréal, Canada.

Senior Research Scientist

Project lead for [PipelineRL](#), an open-source RL framework for LLMs.

Achieved 200% training throughput improvement via sequence packing optimization.

Enabled [Multi-Modal RL](#) with ChartQA visual reasoning implementation.

Replicated [Open Reasoner Zero](#) results while doubling system throughput using novel [inflight weight update](#).

Sept 2019 - Google DeepMind, London, United Kingdom.

Jan 2020 Research Intern

Investigated Bayesian reinforcement learning for exploration in continuous control.

Selected Publications

- 2025 **Piché, A.**, Pardinas, R., Kamaloo, E., Bahdanau, D., *PipelineRL: Faster on-policy Reinforcement Learning for long sequence generation*, Submitted to Neurips 2025.
- 2024 **Piché, A.**, Milios, A., Bahdanau, D., Pal, C., *LLMs can learn self-restraint through iterative self-reflection*, arXiv preprint arXiv:2405.13022.
- 2023 Long, S., **Piché, A.**, Zantedeschi, V., Schuster, T., Drouin, A., *Causal Discovery with Language Models as Imperfect Experts*, Structured Probabilistic Inference & Generative Modeling (SPIGM) workshop at ICML 2023.
- 2023 Long, S., Schuster, T., **Piché, A.**, *Can large language models build causal graphs?*, Causal Machine Learning for Real-World Impact Workshop (CML4Impact) at NeuRIPS 2022.
- 2023 **Piché, A.**, Thomas, V., Marino, J., Pardinas, R., Marconi, G. M., Pal, C., Khan, M. E., *Bridging the Gap Between Target Networks and Functional Regularization*, Transactions on Machine Learning Research, <https://openreview.net/forum?id=BFvoemrmqX>.
- 2021 Marino, J., **Piché, A.**, Ialongo, A. D., Yue, Y., *Iterative amortized policy optimization.*, Advances in Neural Information Processing Systems, 34, 15667-15681..
- 2018 **Piché, A.**, Thomas, V., Ibrahim, C., Bengio, Y., Pal C., *Probabilistic planning with sequential monte carlo methods*, In International Conference on Learning Representations.